



A Multi-thread Sort

Linghan Zhang & Ning Xie

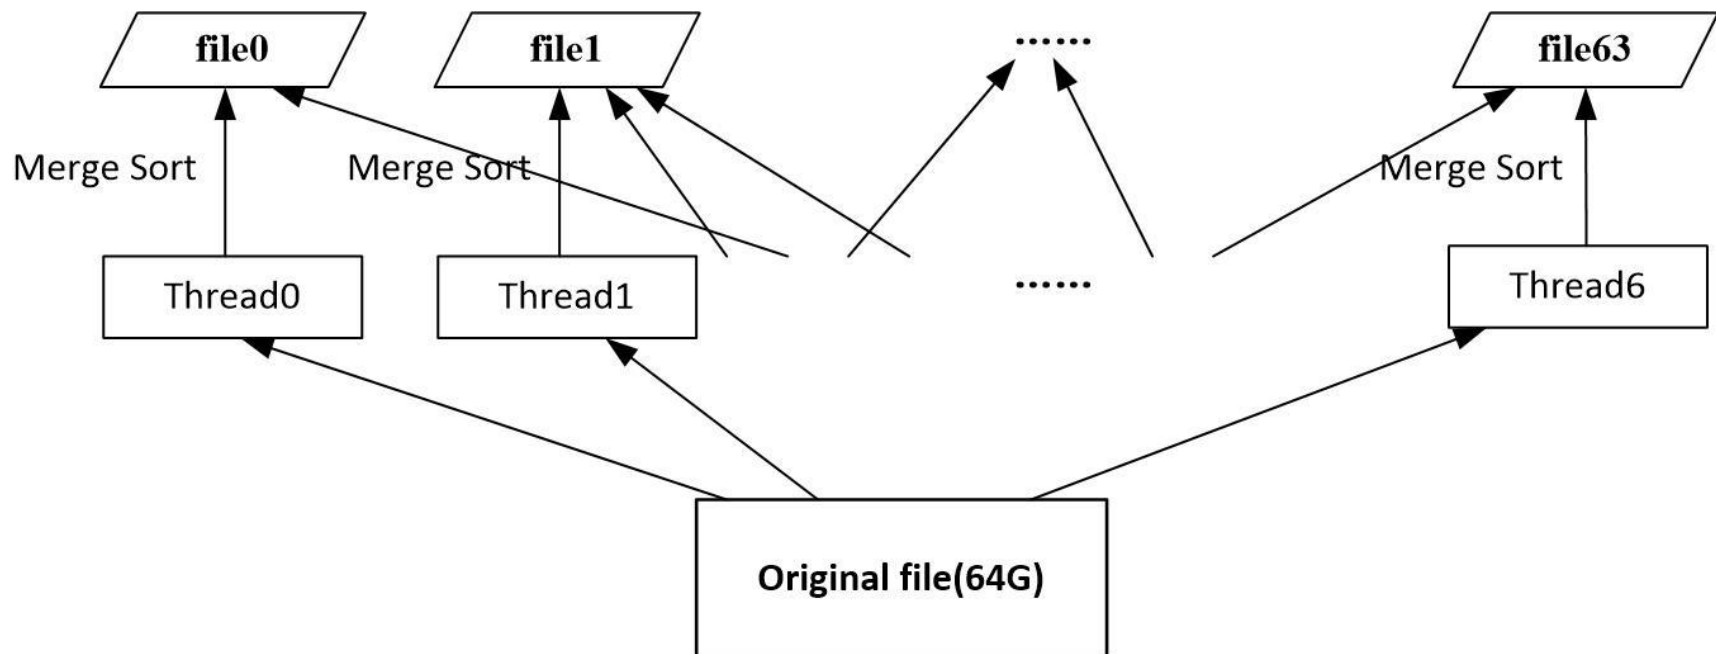


Main Idea

- **Do sorting on a single machine**
 - Avoid data transmission time which might be affected by network
- **Use Merge Sort**
 - Fast and stable

Multi-thread Merge Sort – Phase 1

Divide all data into 64 parts, merge sort and store them into 64 files.



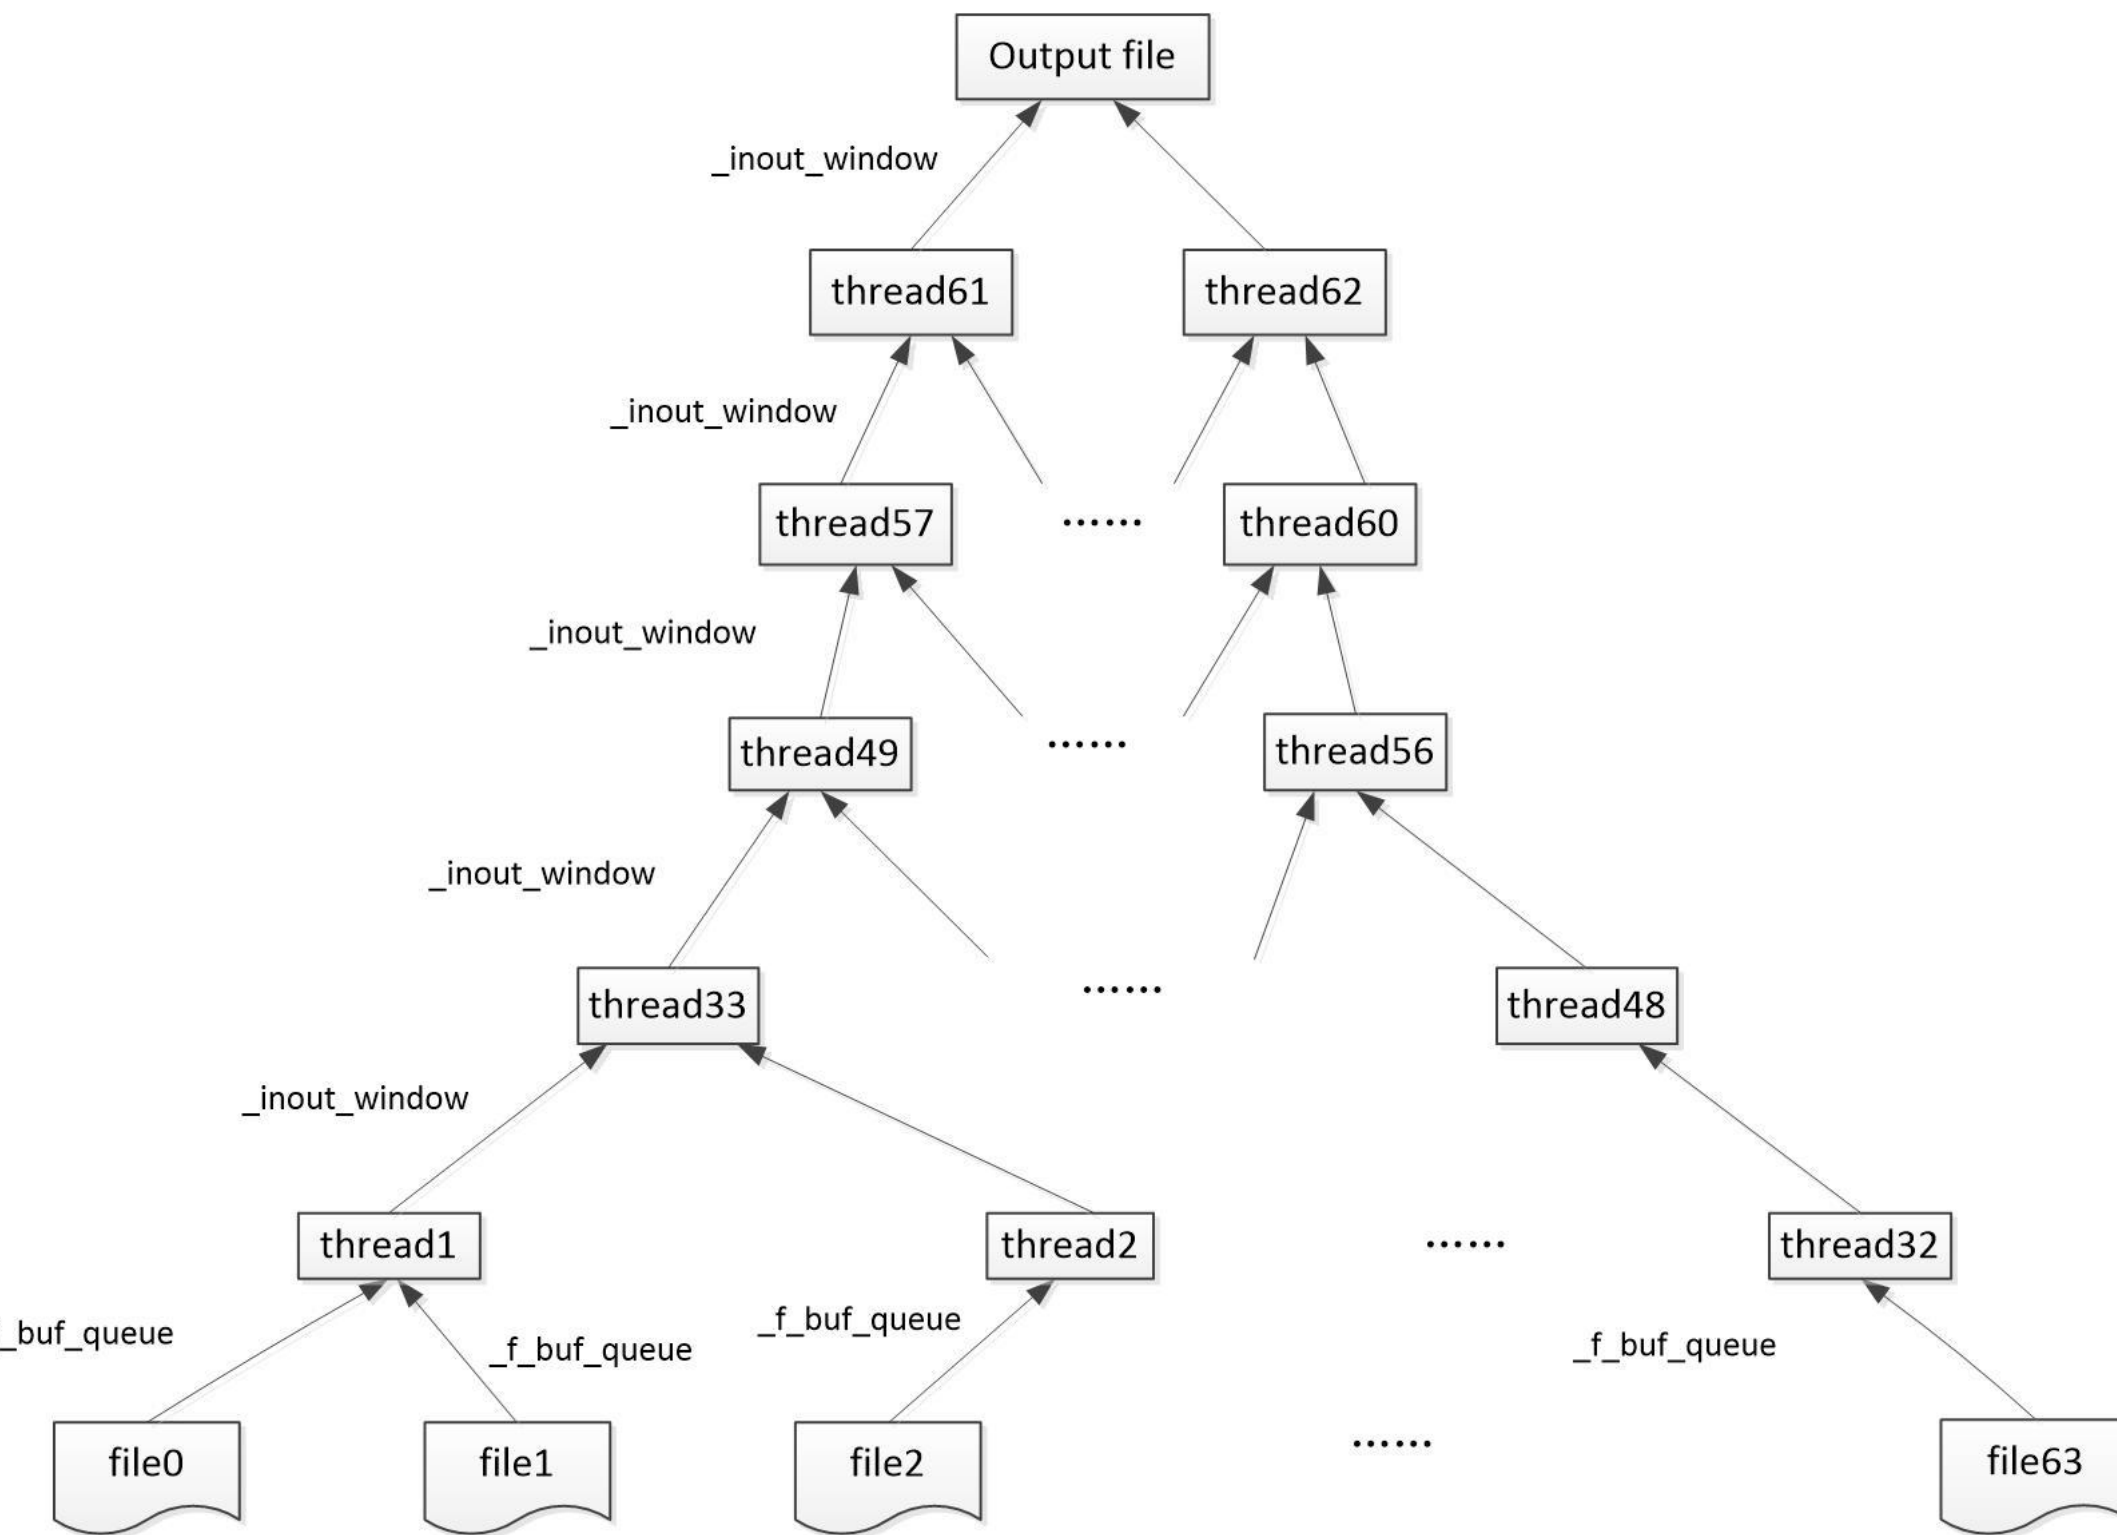
Multi-thread Merge Sort – Phase 2

Create $32+16+8+4+2=62$ threads and merge data

The 62 threads are logically divided into 5 layers:

- Threads of bottom layer(**layer 1**) read from all sorted files in phase 1, and output to window;
- Threads of middle layers(**layers 2-4**) read from lower-layers, merge and output to windows;
- Threads of top layer(**layer 5**) write every 10 number to final output file.





Data Structure Design

```
#define sortqueue_len 2500000 //~20MB, ~2.5M numbers
```

```
typedef struct _inout_window {  
    sort_window * p_in_window1; // read-in window1  
    sort_window * p_in_window2; // read-in window2  
    sort_window * p_out_window; // output window  
    int in_fd1; // file descriptor of input file1 (just use for layer-1 threads)  
    int in_fd2; // file descriptor of input file2 (just use for layer-1 threads)  
    int out_fd; // file descriptor of final output file (just use for layer-5 threads)  
} inout_window;
```

```
typedef struct _sort_window {  
    pthread_mutex_t mutex; // mutex lock  
    pthread_cond_t cond_r; // condition variable for allowing read  
    pthread_cond_t cond_w; // condition variable for allowing write  
    unsigned long long s_queue[sortqueue_len+1]; //circular queue  
    int head; // pointer to the begin of the circular queue  
    int tail; // pointer to the end of the circular queue  
} sort_window;
```

Running Result

Test0: 28min27sec

```
388000000 / 400000000 numbers completed.  
392000000 / 400000000 numbers completed.  
396000000 / 400000000 numbers completed.  
400000000 / 400000000 numbers completed.  
Time used so far is 0 h 42 m 57 s.  
  
real    42m56.727s  
user    20m46.853s  
sys     0m53.756s
```

Test1: 29min13s

```
392000000 / 400000000 numbers completed.  
396000000 / 400000000 numbers completed.  
400000000 / 400000000 numbers completed.  
Time used so far is 0 h 29 m 13 s.  
  
real    29m13.221s  
user    7m39.167s  
sys     0m53.805s
```



Thank you!

